


OPEN-SOURCE MODEL ASSESSMENT

Gemma 3



Google · Released 2025-03-12 · Prepared for Enterprise · Report date 2026-07-16

EXECUTIVE SUMMARY

The Gemma Terms permit commercial use and redistribution of the weights provided the Gemma Terms are attached, but prohibit using the weights to train or improve other models, exceeding 1 million monthly users without a separate Google license, and removing or modifying the Gemma Terms. This restricted/custom license places Gemma 3 at L1 on the OMSF Openness Ladder. For an enterprise reader, DIT recommends: adopt with caution.

OUR RECOMMENDATION

ADOPT WITH CAUTION

OPENNESS GRADE

L1 · Open weights with a restricted/custom license

LICENSE POSTURE

Google license required above 1 million monthly users

DECISION AT A GLANCE

PRIVATE

ADOPT

ENTERPRISE

ADOPT WITH CAUTION

NON-PROFIT

ADOPT WITH CAUTION

01 Openness Ladder verdict

L1 · Open weights with a restricted/custom license

The Gemma Terms permit commercial use and redistribution of the weights provided the Gemma Terms are attached, but prohibit using the weights to train or improve other models, exceeding 1 million monthly users without a separate Google license, and removing or modifying the Gemma Terms. This restricted/custom license places Gemma 3 at L1 on the OMSF Openness Ladder.

PRIMARY SOURCE

[P1] Gemma Terms of Use (P1) <https://ai.google.dev/gemma/terms> · 2025-03-12

02 License summary

YOU ARE ALLOWED TO

- Commercial use within the Gemma Terms
- Redistribution of the weights with the Gemma Terms attached

YOU ARE NOT ALLOWED TO

- Use the weights to train or improve other models
- Exceed 1 million monthly users without a Google license
- Remove or modify the Gemma Terms

COMMERCIAL USE

Commercial use is permitted under the Gemma Terms. If a product or service has more than 1 million monthly users, you must request a license from Google.

SCALE / COMMERCIAL TRIGGER
Google license required above 1 million monthly users

03 Deployment and technical profile

MINIMUM HARDWARE
3 GB VRAM (Gemma 3 4B Q4_K_M)

RUNS FULLY OFFLINE
Yes

DEPLOYMENT VERDICT
runs on a single consumer GPU

TECHNICAL DETAIL

CONTEXT LENGTH
128K tokens

PARAMETERS
1B / 4B / 12B / 27B (dense)

DISTRIBUTION FORMATS
Safetensors, GGUF

PRECISION
BF16

QUANTIZATION IMPACT ON VRAM
For the 27B variant, full BF16 precision requires ~54 GB VRAM; quantizing to GGUF Q4_K_M reduces this to ~14 16 GB ("H4 bit). Similar proportional savings apply to the 12B and 4B variants.

MINIMUM VRAM
3 GB (Gemma 3 4B Q4_K_M)

VRAM BY VARIANT (MIN, QUANTIZED)



MODEL SERIES

Variant	Parameters	Min VRAM	Quant options	Notes
Gemma 3 27B	27B	17 GB (Q4_K_M)	Q4_K_M, Q8_0	Single 24GB GPU
Gemma 3 12B	12B	8 GB (Q4_K_M)	Q4_K_M	Consumer GPU
Gemma 3 4B	4B	3 GB (Q4_K_M)	Q4_K_M	Laptop / edge

04 Decision frames

PRIVATE / INDIVIDUAL

Private individuals can use Gemma 3 under the Gemma Terms for personal projects, research, or prototyping. Commercial use is allowed as long as the resulting service does not exceed 1 million monthly active users, and the weights must not be used to train other models.

ADOPT

ENTERPRISE

Enterprises may deploy Gemma 3 commercially, but must monitor monthly active users; exceeding 1 million MAU triggers a requirement to obtain a separate license from Google. Additionally, the license forbids using the weights to train or improve other models. These obligations necessitate compliance tracking.

ADOPT WITH CAUTION

NON-PROFIT / PUBLIC

Nonprofit organizations enjoy the same commercial use permission as enterprises, subject to the 1 million MAU threshold and the prohibition on using the weights to train other models. If the nonprofit's public-facing application stays below the threshold, usage is unrestricted.

ADOPT WITH CAUTION

05Where sources disagree

None

06Sources and attribution

This report was synthesized with the DIT OpenModel Synthesis Framework (OMSF), a practical, buyer-side adaptation of the Model Openness Framework (LF AI and Data Foundation, 2024; arXiv:2403.13784) and the OSI Open Source AI Definition. Figures are sourced from the web at report date and reflect publicly available information only.

[P1]

1. Gemma 3 on Hugging Face <https://huggingface.co/google/gemma-3-27b-it>

[P1]

2. Gemma 3 official <https://ai.google.dev/gemma>

[P1]

3. Gemma Terms of Use <https://ai.google.dev/gemma/terms>

[P2]

4. Gemma 3 - AI Model Profile - D-Central <https://d-central.tech/ai/model/gemma-3> · 2025-03-12

[P2]

5. Google DeepMind Releases Gemma 3 with Open Weights and Multimodal <https://udit.co/blog/google-deepmind-gemma-3-open-weights-multimodal>

[P1]

6. Mungert/gemma-3-4b-it-gguf <https://huggingface.co/Mungert/gemma-3-4b-it-gguf>

[P3]

7. Which gemma-3 27B for 32GB Vram <https://discuss.vllm.ai/t/which-gemma-3-27b-for-32gb-vram/1625>

[P2]

8. Gemma 3 QAT Models: Bringing state-of-the-Art AI to Consumer GPUs <https://developers.googleblog.com/en/gemma-3-quantized-aware-trained-state-of-the-art-ai-to-consumer-gpus>

[P2]

9. gemma-3 <https://lmstudio.ai/models/gemma-3>