


OPEN-SOURCE MODEL ASSESSMENT

Mistral



L2 · Open weights under a permissive license

Mistral AI · Released 2025-07-22 · Prepared for Private / individual · Report date 2026-07-16

EXECUTIVE SUMMARY

The Apache 2.0 license permits commercial use, modification, and redistribution, requiring only preservation of license and notice, which satisfies the L2 openness rung. For a private / individual reader, DIT recommends: adopt.

OUR RECOMMENDATION

ADOPT

OPENNESS GRADE

L2 · Open weights under a permissive license

LICENSE POSTURE

Commercial use permitted

DECISION AT A GLANCE

PRIVATE

ADOPT

ENTERPRISE

ADOPT

NON-PROFIT

ADOPT

01 Openness Ladder verdict

L2 · Open weights under a permissive license

The Apache 2.0 license permits commercial use, modification, and redistribution, requiring only preservation of license and notice, which satisfies the L2 openness rung.

PRIMARY SOURCE

[P1] Apache 2.0 (open-weight releases) (P1) <https://mistral.ai/licenses> · 2025-07-22

02 License summary

YOU ARE ALLOWED TO

- Commercial use
- Modification
- Redistribution

YOU ARE NOT ALLOWED TO

- Remove the Apache 2.0 notice
- Use the Mistral trademarks without permission

COMMERCIAL USE

Mistral open-weight releases are under Apache 2.0, which permits commercial use, modification and redistribution with attribution.

03 Deployment and technical profile

MINIMUM HARDWARE

5 GB VRAM (Q4_K_M) for the 7B variant

RUNS FULLY OFFLINE

Yes

THROUGHPUT (P2)

~25-30 tokens/s on GTX 1070 (8GB VRAM) with Q6_K quantization

DEPLOYMENT VERDICT

runs on a single consumer GPU

TECHNICAL DETAIL

CONTEXT LENGTH

128K tokens (Mistral Large 3)

PARAMETERS

Mistral Large 3: ~100B+ (per model card); 7B / 8B open-weight also available

DISTRIBUTION FORMATS

Safetensors, GGUF

PRECISION

BF16

QUANTIZATION IMPACT ON VRAM

BF16 7B ~14 GB VRAM; Q4_K_M GGUF ~5 GB VRAM ("H65% reduction)

TOKENS / SEC (P2)

~25-30 tokens/s on GTX 1070 (8GB VRAM) with Q6_K quantization

MINIMUM VRAM

5 GB

VRAM BY VARIANT (MIN, QUANTIZED)

Mistral 7B

5 GB (Q4_K_M)

MODEL SERIES

Variant	Parameters	Min VRAM	Quant options	Notes
Mistral 7B	7B	5 GB (Q4_K_M)	Q4_K_M, Q8_0	Consumer GPU
Mistral Large 3	~100B+	Not available	GGUF	Check model card for VRAM

04 Decision frames

PRIVATE / INDIVIDUAL Apache 2.0 license allows commercial use, modification, and redistribution; the 7B variant runs on consumer hardware with as little as 5 GB VRAM when quantized.	ADOPT
ENTERPRISE No usage based thresholds or royalties apply; enterprises can deploy the model commercially under Apache 2.0 with only attribution required.	ADOPT
NON-PROFIT / PUBLIC Same permissive terms apply; nonprofits can freely use, modify, and redistribute the model under Apache 2.0.	ADOPT

05 Where sources disagree

Source [5] reports 6 GB VRAM for Mistral 7B (likely FP16) whereas authoritative facts list 5 GB VRAM for Q4_K_M quantization; source [4] mentions 16 32 GB RAM or "e16 GB VRAM GPU for Mistral Small 4 Q4/Q5, which pertains to a different model and not the Mistral 7B/Large 3 family.

06 Sources and attribution

This report was synthesized with the DIT OpenModel Synthesis Framework (OMSF), a practical, buyer-side adaptation of the Model Openness Framework (LF AI and Data Foundation, 2024; arXiv:2403.13784) and the OSI Open Source AI Definition. Figures are sourced from the web at report date and reflect publicly available information only.

[P1] 1. Mistral Licenses <https://mistral.ai/licenses>

[P1] 2. Hardware specs for GGUF 7B/13B/30B parameter models #3847 <https://github.com/ggml-org/llama.cpp/discussions/3847> · Oct 29, 2023