


OPEN-SOURCE MODEL ASSESSMENT

Qwen3



L2 · Open weights under a permissive license

Alibaba (Qwen Team) · Released 2025-04-29 · Prepared for Private / individual · Report date 2026-07-16

EXECUTIVE SUMMARY

The license is explicitly identified as 'Apache License Version 2.0', a permissive license that permits commercial use, modification, and distribution provided attribution and license notice are retained. For a private / individual reader, DIT recommends: adopt.

OUR RECOMMENDATION

ADOPT

OPENNESS GRADE

L2 · Open weights under a permissive license

LICENSE POSTURE

Commercial use permitted

DECISION AT A GLANCE

PRIVATE

ADOPT

ENTERPRISE

ADOPT

NON-PROFIT

ADOPT

01 Openness Ladder verdict

L2 · Open weights under a permissive license

The license is explicitly identified as 'Apache License Version 2.0', a permissive license that permits commercial use, modification, and distribution provided attribution and license notice are retained.

PRIMARY SOURCE

[P1] Apache 2.0 (P1) <https://www.apache.org/licenses/LICENSE-2.0> · 2025-04-29

02 License summary

YOU ARE ALLOWED TO

- Commercial use
- Modification
- Redistribution
- Fine-tuning and derivative works

YOU ARE NOT ALLOWED TO

- Remove the Apache 2.0 license or attribution notices
- Use the Qwen trademarks without permission

COMMERCIAL USE

Apache 2.0 permits commercial use, modification and redistribution provided attribution and the license notice are retained.

03 Deployment and technical profile

MINIMUM HARDWARE

20 GB VRAM (Q4_K_M) for Qwen3-32B dense variant

RUNS FULLY OFFLINE

Yes

DEPLOYMENT VERDICT

runs on a single consumer GPU

TECHNICAL DETAIL

CONTEXT LENGTH

32K tokens native (up to 128K with YaRN)

PARAMETERS

Dense 0.6B to 32B; MoE 235B-A22B (22B active)

DISTRIBUTION FORMATS

Safetensors, GGUF

PRECISION

BF16

QUANTIZATION IMPACT ON VRAM

Moving from BF16 ("H2 bytes/param) to GGUF Q4_K_M ("H0.5 bytes/param) reduces VRAM roughly 4 fold; e.g., Qwen3-32B drops from ~64 GB to ~20 GB, and the MoE 235B A22B drops from ~470 GB to ~130 GB.

MINIMUM VRAM

20 GB (Q4_K_M, Qwen3-32B)

VRAM BY VARIANT (MIN, QUANTIZED)



MODEL SERIES

Variant	Parameters	Min VRAM	Quant options	Notes
Qwen3-32B	32B dense	20 GB (Q4_K_M)	Q4_K_M, Q8_0	Single 24GB GPU
Qwen3-235B-A22B	235B total / 22B active (MoE)	130 GB (Q4_K_M)	Q4_K_M	Multi-GPU

04 Decision frames

PRIVATE / INDIVIDUAL

Qwen3 is released under the permissive Apache 2.0 license, allowing commercial use, modification, and redistribution. The dense 32B variant fits on a single consumer grade 24 GB GPU when quantized to Q4_K_M, enabling offline operation without licensing friction.

ADOPT

ENTERPRISE

Enterprises can freely deploy Qwen3 commercially under Apache 2.0 with no usage based fees or MAU thresholds. Larger MoE variants require multi GPU servers but remain free of licensing restrictions; attribution and preservation of the license notice are the only obligations.

ADOPT

NON-PROFIT / PUBLIC

Non profit organizations benefit from the same permissive terms: no cost, no usage limits, and the ability to fine tune or derivative works. The model runs offline on modest hardware when quantized, supporting mission driven applications.

ADOPT

05 Where sources disagree

Source 5 (D Central) references a Qwen3 30B A3B variant that is not listed in the authoritative facts for the Qwen3 base release; this claim relies on a single P3 source and is unconfirmed for this report.

06 Sources and attribution

This report was synthesized with the DIT OpenModel Synthesis Framework (OMSF), a practical, buyer-side adaptation of the Model Openness Framework (LF AI and Data Foundation, 2024; arXiv:2403.13784) and the OSI Open Source AI Definition. Figures are sourced from the web at report date and reflect publicly available information only.

[P1] 1. Qwen3 blog <https://qwenlm.github.io/blog/qwen3/>

[P1] 2. Qwen3 on Hugging Face <https://huggingface.co/Qwen>

[P1] 3. Mungert/Qwen3-8B-GGUF · Hugging Face <https://huggingface.co/Mungert/Qwen3-8B-GGUF>

[P3] 4. Qwen 3 - AI Model Profile - D-Central <https://d-central.tech/ai/model/qwen-3>

[P3] 5. How to choose a Qwen3-8B quantization method: deciding AWQ, GPTQ, FP8, and GGUF by use <https://tomodahinata.com/en/blog/qwen3-quantization-awq-gptq-fp8-gguf-comparison-guide>